



MANGALORE UNIVERSITY

**MASTER OF SCIENCE
IN
COMPUTER SCIENCE**

**22CSP309: Artificial Intelligence
and
Machine Learning Lab**

**SUBMITTED
BY**

**III SEMESTER MSC
Computer Science Students**

**SUBMITTED
TO**

**Chairperson
The Department of Computer Science**

Examiner:

1.

2.

Mangalore University
Dept. of Post-Graduate Studies and Research in Computer Science
Mangalagangothri – 574199

INDEX

SL.NO	CONTENT	Page. No
1	Implement the Principle Component Analysis (PCA) Algorithm.	3-5
2	Write a program to perform Linear Regression with-out built-in function on the dataset	6-7
3	Write a program to perform Linear Regression withbuilt-in function on the dataset	8-10
4	Implement Support Vector Machine(SVM) Classifier with built-in function on the Suitable dataset.	11-15
5	Implement Gradient Descent algorithm for the single variable.	16-19
6	Implement Text Preprocessing 1)Removing punctuations 2)Removing digits 3)Removing Special characters 4)Tokenization 5)Stemminig 6)Lemmatization 7)Removal of Stopwords 8)Demojize	20-23
7	Implement feature extraction using TF-IDF of the models(SVM, KNN, Logistic Regression, Decision tree)	24-29
8	Write a program to perform Decision Tree classifier with built-in function on the dataset	30-32
9	Write a program to performDensity Based Special Clustering with built-in function on the dataset	33-35
10	Implement K-Means Clustering with built-in function on the dataset	36-38
11	Implement K-Nearest Neighbour classifier with built-in function on the dataset	39-42

1. Write A Python Code To Demonstrate Principal Components Analysis

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

```
# Step 1: Create some sample data (replace this with your dataset)
data = pd.read_csv('/content/drive/MyDrive/ex1data1.txt') # 100 samples with 3
features
np.set_printoptions(precision=4, suppress=True)
print("Formatted Array:")
print(data)
```

Formatted Array:

```
6.1101 17.592
0 5.5277 9.13020
1 8.5186 13.66200
2 7.0032 11.85400
3 5.8598 6.82330
4 8.3829 11.88600
.. ...
91 5.8707 7.20290
92 5.3054 1.98690
93 8.2934 0.14454
94 13.3940 9.05510
95 5.4369 0.61705
```

[96 rows x 2 columns]

```
# Step 2: Standardize the data
mean = np.mean(data, axis=0)
#std_dev = np.std(data, axis=0)
standardized_data = (data - mean)
print(mean)
print(standardized_data)
```

```
6.1101 8.181151
17.592 5.716709
6.1101 17.592
0 -2.653451 3.413491
1 0.337449 7.945291
```

```
2 -1.177951 6.137291
3 -2.321351 1.106591
4 0.201749 6.169291
.. ... ..
91 -2.310451 1.486191
92 -2.875751 -3.729809
93 0.112249 -5.572169
94 5.212849 3.338391
95 -2.744251 -5.099659
```

```
[96 rows x 2 columns]
```

```
# Step 3: Compute the covariance matrix
covariance_matrix = np.cov(standardized_data, rowvar=False)
size_cc = covariance_matrix.size
shape_cc = covariance_matrix.shape
print (size_cc, shape_cc)
print(covariance_matrix)
```

```
4 (2, 2)
[[15.089 18.3112]
 [18.3112 29.2135]]
```

```
# Step 4: Compute the eigenvalues and eigenvectors of the covariance matrix
eigenvalues, eigenvectors = np.linalg.eigh(covariance_matrix)
print(eigenvalues)
print(eigenvectors)
```

```
[ 2.5254 41.7771]
[[-0.8246 0.5658]
 [ 0.5658 0.8246]]
```

```
# Step 5: Sort eigenvalues and corresponding eigenvectors in descending order
sorted_indices = np.argsort(eigenvalues)[::-1]
eigenvalues = eigenvalues[sorted_indices]
eigenvectors = eigenvectors[:, sorted_indices]
print(eigenvalues)
print(eigenvectors)
```

```
[41.7771 2.5254]
[[ 0.5658 -0.8246]
 [ 0.8246 0.5658]]
```

Step 6: Choose the number of components (or a threshold for explained variance)

n_components = 1 # Choose the number of principal components

Step 7: Select the top 'n_components' eigenvectors

selected_eigenvectors = eigenvectors[:, :n_components]

Step 8: Project the data onto the selected eigenvectors to obtain the principal components

final_result = np.dot(standardized_data, selected_eigenvectors)

Step 9: Print the final result

print("Final Result after PCA:")

print(final_result)

Final Result after PCA:

```
[[ 1.3135]
 [ 6.7424]
 [ 4.3942]
 [-0.4008]
 [ 5.2012]
 [-1.5271]
 [ 5.4056]
 [-0.2317]
 [-3.3356]
 [-3.4298]
 [11.456 ]
 [-3.4967]
 [ 1.3729]
 [-5.5606]
 [-3.4023]
 [-1.3669]
 [-5.9776]
 [-2.6936]
 [-0.8982]
 [-3.2529]
 [20.074 ]
 [-2.7212]
 [-1.4859]
 [-3.6523]
 [20.0425]
```

2. WRITE A LINEAR REGRESSION PROGRAM TO PERFORM WITHOUT BUILT-IN FUNCTION.

```
from google.colab import drive
drive.mount('/content/drive')

import pandas as pd
df=pd.read_csv('/content/drive/MyDrive/Salary_Data.csv')
df
```

output:

```
   x   y
0  1   2
1  2   4
2  3   5
3  4   4
4  5   5
```

```
import matplotlib.pyplot as plt
mean_x=df['YearsExperience'].mean()
mean_y=df['Salary'].mean()
print(mean_x)
print(mean_y)
```

output:

```
3.0
4.0
```

```
numerator = 0
denominator = 0
for i in range(len(df['YearsExperience'])):
    numerator += (df['YearsExperience'][i] - mean_x) * (df['Salary'][i] - mean_y)
    denominator += (df['YearsExperience'][i] - mean_x) ** 2
m= numerator/denominator
```

```
c=mean_y-m*mean_x
print('slope:',m)
print('intercept:',c)
```

output:
slope: 0.6 intercept: 2.2

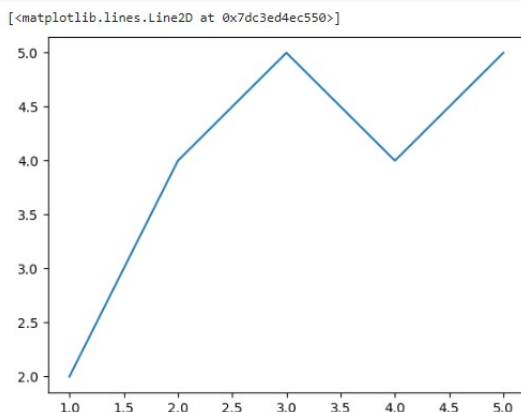
```
predicted_y = [(m * xi) + c for xi in df['YearsExperience']]
print('predicty:',predicted_y)
```

output:
predicty: [2.8000000000000003, 3.4000000000000004, 4.0, 4.6, 5.2]

```
yr = [df['Salary'][i] - predicted_y[i] for i in range(len(df))]
print(yr)
```

output:
[-0.8000000000000003, 0.5999999999999996, 1.0, -0.5999999999999996, -0.20000000000000018]

```
import numpy as np
R=np.mean(yr)
print(R)
output:
-8.881784197001253e-17
x=df['YearsExperience']
y=df['Salary']
plt.plot(x,y)
```



3. Linear regression program with built-in function.

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt # data plotting

dataset = pd.read_csv('/content/drive/MyDrive/AI and ML/ex1data1.txt')
dataset.head()
```

```
      6.1101   17.592
0    5.5277   9.1302
1    8.5186  13.6620
2    7.0032  11.8540
3    5.8598   6.8233
4    8.3829  11.8860
```

```
X = dataset.iloc[:, :-1].values #independent variable array
y = dataset.iloc[:, 1].values #dependent variable vector
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test =
train_test_split(X, y, test_size=1/3, random_state=0)
```

```
from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(X_train, y_train)
```

```
▼ LinearRegression
LinearRegression()
```

```
y_pred = regressor.predict(X_test)
y_pred
```



```
array([ 8.94195531, 6.10107029, 4.33817991, 3.48933772, 19.86861533,
       3.58423571, 2.15715616, 2.13107377, 3.58796177, 8.10242826,
       3.55757112, 2.07483361, 3.17728051, 2.75169502, 3.73618895,
       3.85635426, 3.54976969, 18.24312328, 5.74464972, 2.06924452,
       5.29472842, 13.54247042, 3.00681343, 3.17239006, 2.08740905,
       2.63630371, 6.17198179, 2.51439181, 2.22504025, 4.9615957 ,
       2.21444428, 2.85474376])
```

y_test

```
array([ 7.0467 , 4.2415 , 11.854 , 2.4756 , 20.992 , 0.67861,
       0.56077, 2.0576 , -1.4211 , 7.7754 , 1.4233 , 2.8214 ,
       1.2784 , 0.71618, 6.5987 , 5.9966 , 5.1875 , 22.638 ,
       6.5426 , 3.8166 , 6.7318 , 12.054 , 6.8233 , 0.92695,
       5.1337 , 1.0179 , 12. , 0.61705, 0.20421, 3.8845 ,
       -0.74279, 0.47953])
```

#plot for the TRAIN

```
plt.scatter(X_train, y_train, color='red') # plotting the observation line
plt.plot(X_train, regressor.predict(X_train), color='blue') # plotting the
regression line
plt.title("Salary vs Experience (Training set)") # stating the title of the graph
plt.xlabel("Years of experience") # adding the name of x-axis
plt.ylabel("Salaries") # adding the name of y-axis
plt.show() # specifies end of graph
```



```
#plot for the TEST
plt.scatter(X_test, y_test, color='red')
plt.plot(X_train, regressor.predict(X_train), color='blue') # plotting the
regression line
plt.title("Salary vs Experience (Testing set)")
plt.xlabel("Years of experience")
plt.ylabel("Salaries")
plt.show()
```



```
print("Coefficient:",regressor.coef_)
print("Intercept",regressor.intercept_)
```

Coefficient: [1.16439258]
Intercept -3.81629423778271

4. Implement Support Vector Machine (Svm) Classifier With Suitable Datasets

```
# Support Vector Machine (SVM)

# Importing the libraries

import numpy as np

import matplotlib.pyplot as plt

import pandas as pd

# Importing the dataset

dataset = pd.read_csv('/content/Social_Network_Ads.csv')

dataset
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1
400 rows × 5 columns					

```

X = dataset.iloc[:, [2, 3]].values

y = dataset.iloc[:, 4].values

# Splitting the dataset into training and test set.

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25,
random_state = 0)


# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)


# Fitting SVM to the Training set
from sklearn.svm import SVC
classifier = SVC(kernel = 'linear', random_state = 0)
classifier.fit(X_train, y_train)
: SVC
SVC(kernel='linear', random_state=0)


# Predicting the Test set results
y_pred = classifier.predict(X_test)
y_pred

: array([0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1,
        0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
        1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1,
        0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1,
        0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1])


```

```
# Making the Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
cm

:array([[66, 2],
       [ 8, 24]])
```

```
# Visualizing the Training set results
```

```
from matplotlib.colors import ListedColormap
# Create a meshgrid to plot the decision boundary
X1, X2 = np.meshgrid(np.arange(start=X_train[:, 0].min() - 1, stop=X_train[:,
0].max() + 1, step=0.01), np.arange(start=X_train[:, 1].min() - 1, stop=X_train[:,
1].max() + 1, step=0.01)
# Use the classifier to predict the class labels for each point in the meshgrid
Z = classifier.predict(np.array([X1.ravel(), X2.ravel()]).T)
Z = Z.reshape(X1.shape)
# Create a color map for the plot
cmap = ListedColormap(['red', 'green'])
# Plot the training set data points
plt.figure(figsize=(8, 6))
plt.contourf(X1, X2, Z, alpha=0.5, cmap=cmap)
plt.scatter(X_train[y_train == 0, 0], X_train[y_train == 0, 1], color='red',
label='Not Purchased')
plt.scatter(X_train[y_train == 1, 0], X_train[y_train == 1, 1], color='green',
label='Purchased')
plt.title('Decision Tree Classifier - Training Set')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()
```



```
# Visualizing the Testing set results
from matplotlib.colors import ListedColormap
# Create a meshgrid of feature values
X1, X2 = np.meshgrid(np.arange(start = X_test[:, 0].min() - 1, stop = X_test[:,
0].max() + 1, step = 0.01),np.arange(start = X_test[:, 1].min() - 1, stop = X_test[:,
1].max() + 1, step = 0.01))
# Use the trained classifier to make predictions on the meshgrid points
Z = classifier.predict(np.array([X1.ravel(), X2.ravel()]).T)
Z = Z.reshape(X1.shape)
# Create a colormap for the two classes
cmap = ListedColormap(['red', 'green'])
# Plot the contour filled by the predictions
plt.contourf(X1, X2, Z, alpha = 0.5, cmap = cmap)
# Scatter plot the actual data points
plt.scatter(X_test[y_test == 0, 0], X_test[y_test == 0, 1], color = 'red', label = 'Not
Purchased')
plt.scatter(X_test[y_test == 1, 0], X_test[y_test == 1, 1], color = 'green', label =
'Purchased')
# Add labels and legend
plt.title('SVM - Testing Set')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
# Show the plot
plt.show()
```



```
from sklearn.metrics import accuracy_score  
print ('Accuracy : ', accuracy_score(y_test, y_pred))  
: Accuracy : 0.9
```


5. Implement Gradient descent algorithm for the single variable

```
from google.colab import drive
drive.mount('/content/drive')
```

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
# %matplotlib inline
```

```
df = pd.read_csv('/content/drive/MyDrive/Copy of ex1data1.txt', sep=',', header=None)
df.columns = ['population', 'profit']
df
```

```
ax = sns.scatterplot(x='population', y='profit', data=df)
ax.set(xlabel='Population of City in 10,000s', ylabel='Profit in $10,000s', title='Scatter plot of training data')
```



```
m = df.shape[0]
X = np.hstack((np.ones((m,1)), df.population.values.reshape(-1,1)))
y = np.array(df.profit.values).reshape(-1,1)
theta = np.zeros(shape=(X.shape[1],1))
```

```
iterations = 1500
alpha = 0.01
print(m, X)
```

```
def compute_cost_one_variable(X, y, theta):
    m = y.shape[0]
    h = X.dot(theta)
```



```

J = (1/(2*m)) * (np.sum((h - y)**2))
return J

J = compute_cost_one_variable(X, y, theta)
print('With theta = [0 ; 0]\nCost computed =', J)
print('Expected cost value (approx) 32.07')

J = compute_cost_one_variable(X, y, [[-1],[2]])
print('With theta = [-1 ; 2]\nCost computed =', J)
print('Expected cost value (approx) 54.24')

def gradient_descent(X, y, theta, alpha, num_iters):
    m = y.shape[0]
    J_history = np.zeros(shape=(num_iters, 1))

    for i in range(0, num_iters):
        h = X.dot(theta)
        diff_hy = h - y


        delta = (1/m) * (diff_hy.T.dot(X))
        theta = theta - (alpha * delta.T)
        J_history[i] = compute_cost_one_variable(X, y, theta)

    return theta, J_history

theta, _ = gradient_descent(X, y, theta, alpha, iterations)
print('Theta found by gradient descent:\n', theta)
print('Expected theta values (approx)\n -3.6303\n 1.1664')

ax = sns.scatterplot(x='population', y='profit', data=df)
plt.plot(X[:,1], X.dot(theta), color='r')
ax.set(xlabel='Population of City in 10,000s', ylabel='Profit in $10,000s', title='Training
data with linear regression fit')

```



```
y_pred = np.array([1, 3.5]).dot(theta)
f'For population = 35,000, we predict a profit of {y_pred[0]*10000}'
```


```
y_pred = np.array([1, 7]).dot(theta)
f'For population = 70,000, we predict a profit of {y_pred[0]*10000}'
```

```
theta0_vals = np.linspace(-10, 10, 100)
theta1_vals = np.linspace(-1, 4, 100)
```

```
J_vals = np.zeros(shape=(len(theta0_vals), len(theta1_vals)))
```

```
for i in range(0, len(theta0_vals)):
    for j in range(0, len(theta1_vals)):
        J_vals[i,j] = compute_cost_one_variable(X, y, [[theta0_vals[i]], [theta1_vals[j]]])
```

```
ax = plt.contour(theta0_vals, theta1_vals, np.transpose(J_vals), levels=np.logspace(-2,3,20))
plt.plot(theta[0,0], theta[1,0], marker='x', color='r')
plt.xlabel(r'$\theta_0$')
plt.ylabel(r'$\theta_1$')
plt.title('Contour, showing minimum')
```



```
from sklearn.linear_model import LinearRegression
lin_reg = LinearRegression()
lin_reg.fit(df.population.values.reshape(-1,1),
            df.profit.values.reshape(-1,1))

print(lin_reg.intercept_, lin_reg.coef_)
```

6. Implement Text Preprocessing

- 1. Removing punctuations**
- 2. Removing digits**
- 3. Removing Special characters**
- 4. Tokenization**
- 5. Lemmatization**
- 6. Removal of Stopwords**
- 7. Demojize**

```
from google.colab import drive
drive.mount('/content/drive')
```

```
!pip install contractions
```

```
!pip install emoji
```

```
!pip install nltk
```

```
import pandas as pd
import numpy as np
from nltk.tokenize import word_tokenize
import statistics
import nltk
nltk.download('punkt')
nltk.download('stopwords')
nltk.download('whitespace')
nltk.download('wordnet')
import json
from nltk import word_tokenize
from nltk.corpus import stopwords
import string
import contractions
from nltk.stem import WordNetLemmatizer
```

```

from nltk.stem import PorterStemmer
import re
from nltk.tokenize import TweetTokenizer
import emoji
import regex
from nltk.corpus import stopwords
from nltk.stem import SnowballStemmer

```

#read dataset

```

dft=pd.read_csv("/content/drive/MyDrive/principle of data science lab/sample_text
- Sheet1.csv")

```

dft

	tweet_id	text	task1
0	1123757263427186690	hate wen females hit ah nigga with tht bro ☐☐,...	HOF
1	1123733301397733380	RT @airjunebug: When you're from the Bay but y...	HOF
2	1123734094108659712	RT @DonaldJTrumpJr: Dear Democrats: The Americ...	
	NOT		
3	1126951188170199049	RT @SheLoveTimothy: He ain't on drugs he just ...	HOF
4	1126863510447710208	RT @TavianJordan: Summer '19 I'm coming for yo...	
	NOT		
...
995	1126798721025544193	RT @prodnose: Good morning, everyone.\nFollowi...	
	NOT		
996	1126833089190219777	@cheezitking123 this what you get for tryna ge...	NOT
997	1130037092845670400	earphones ko ☐☐☐☐☐☐☐	NOT
998	1127028455651123201	RT @nj_linguist: @realgonegirl @elivalley I th...	NOT
999	1130285076858789889	i'm tired as fuck. and man, physically ain't S...	HOF

1000 rows × 3 columns

#tokenization

```

dft['Text_Tokenized'] = dft['text'].str.lower().apply(word_tokenize)

```

dft

	tweet_id	text	task1	Text_Tokenized
0	1123757263427186690	hate wen females hit ah nigga with tht bro □□,... [hate, wen, females, hit, ah, nigga, with, tht...		HOF
1	1123733301397733380	RT @airjunebug: When you're from the Bay but y... [rt,		HOF
2	1123734094108659712	RT @DonaldJTrumpJr: Dear Democrats: The Americ... NOT [rt,		
3	1126951188170199049	RT @SheLoveTimothy: He ain't on drugs he just ... [rt,		HOF
4	1126863510447710208	RT @TavianJordan: Summer '19 I'm coming for yo... NOT [rt,		

#Preprocessing

```

ps =PorterStemmer()
lemmatiser = WordNetLemmatizer()
english_stopwords = stopwords.words('english')
exclude = set(string.punctuation)
text=dft["text"].iloc[:40]
def preprocess(text):
    text=contractions.fix(text.lower(),slang=True)
    text =re.sub("@ ?[A-Za-z0-9_]+", "", text)
    text= re.sub(r'\d+', "", text)
    text=re.sub(r'$', "", text)
    text= re.sub(r'","', "", text )
    text=re.sub('<.*?>', "",text)
    text=re.sub(r'http\S+', "", text)
    tokens = word_tokenize(text)
    tokens = [lemmatiser.lemmatize(t) for t in tokens]
    tokens = [t for t in tokens if t not in english_stopwords]
    return text

```

```
dft['preprocessed_text'] = dft['text'].apply(preprocess)  
dft['preprocessed_text']
```

```
0    hate wen females hit ah nigga with tht bro ☐☐,...  
1    rt : when you are from the bay but you are rea...  
2    rt : dear democrats: the american people are n...  
3    rt : he are not on drugs he just bored. i be d...  
4    rt : summer ' i am coming for you ! no boring ...
```

```
import emoji
```

```
def remove_emojis(text):  
    temp = emoji.demojize(text,delimiters=("", ""))  
    temp = temp.replace("_", "")  
    return temp
```

```
dft['clean'] = dft['preprocessed_text'].apply(lambda x: remove_emojis(x))  
dft['clean']=dft['clean'].apply(lambda X: preprocess(X))  
dft['clean']
```

```
0    hate wen females hit ah nigga with tht bro fac...  
1    rt : when you are from the bay but you are rea...  
2    rt : dear democrats: the american people are n...  
3    rt : he are not on drugs he just bored. i be d...  
4    rt : summer ' i am coming for you ! no boring ...
```

7. Implement feature extraction using TF-IDF of the models(SVM, KNN, Logistic Regression, Decision tree)

```
from google.colab import drive
drive.mount('/content/drive')
```

```
!pip install emoji
import pandas as pd
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import ENGLISH_STOP_WORDS
import string
import emoji
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import PorterStemmer, WordNetLemmatizer
import re
import nltk
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
```

```
nltk.download('punkt')
nltk.download('stopwords')
nltk.download('wordnet')
```

Load your CSV file into a pandas DataFrame

```
csv_file_path = '/content/drive/MyDrive/lab programs/emotion-labels-test.csv' #
```

Replace with the path to your CSV file

```
df = pd.read_csv(csv_file_path)
df
```


	text	label
0	You must be knowing #blithe means (adj.) Happ...	joy
1	Old saying 'A #smile shared is one gained for ...	joy
2	Bridget Jones' Baby was bloody hilarious 🤪 #Br...	joy
3	@Elaminova sparkling water makes your life spa...	joy
4	I'm tired of everybody telling me to chill out...	joy
...
3137	Why does Candice constantly pout #GBBO 🍷 😞	sadness
3138	@redBus_in #unhappy with #redbus CC, when I ta...	sadness
3139	@AceOperative789 no pull him afew weeks ago, s...	sadness
3140	I'm buying art supplies and I'm debating how s...	sadness
3141	@sainsburys Could you ask your Chafford Hundre...	sadness

#Pre-Processing

ps = PorterStemmer()

lemmatiser = WordNetLemmatizer()

english_stopwords = set(stopwords.words('english'))

exclude = set(string.punctuation)

def preprocess(text):

Convert to lowercase

text = text.lower()

Remove digits

text = re.sub(r'\d+', "", text)

Remove HTML tags

text = re.sub('<.*?>', "", text)

Remove URLs

text = re.sub(r'http\S+', "", text)

Remove non-ASCII characters

text = text.encode("ascii", "ignore").decode()

Remove punctuation

```
text = "".join(ch for ch in text if ch not in exclude)
```

```
# Tokenize, lemmatize, and stem
```

```
tokens = word_tokenize(text)
```

```
tokens = [lemmatiser.lemmatize(t) for t in tokens]
```

```
tokens = [ps.stem(t) for t in tokens]
```

```
# Remove stopwords
```

```
tokens = [t for t in tokens if t not in english_stopwords]
```

```
# Join tokens back into text
```

```
text = " ".join(tokens)
```

```
# Demojize and replace underscores
```

```
text = emoji.demojize(text, delimiters=(" ", " "))
```

```
text = text.replace("_", " ")
```

```
return text
```

```
df['processed_text'] = df['text'].astype(str).apply(preprocess)
```

```
df
```

	text	label
0	You must be knowing #blithe means (adj.) Happ...	joy
1	Old saying 'A #smile shared is one gained for ...	joy
2	Bridget Jones' Baby was bloody hilarious 😂 #Br...	joy
3	@Elaminova sparkling water makes your life spa...	joy
4	I'm tired of everybody telling me to chill out...	joy
...
3137	Why does Candice constantly pout #GBBO 🗨️ 😞	sadness
3138	@redBus_in #unhappy with #redbus CC, when I ta...	sadness
3139	@AceOperative789 no pull him afew weeks ago, s...	sadness
3140	I'm buying art supplies and I'm debating how s...	sadness
3141	@sainsburys Could you ask your Chafford Hundre...	sadness

```
X = df['processed_text']
y = df['label']
```

```
# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```

```
# TF-IDF Vectorization
vectorizer = TfidfVectorizer()
X_train_vec = vectorizer.fit_transform(X_train)
X_test_vec = vectorizer.transform(X_test)
```

```
svm_model = SVC(kernel='linear')
svm_model.fit(X_train_vec, y_train)
```

```
# Decision Tree
dt_model = DecisionTreeClassifier()
dt_model.fit(X_train_vec, y_train)
```

```
# k-Nearest Neighbors (KNN)
knn_model = KNeighborsClassifier()
knn_model.fit(X_train_vec, y_train)
```

```
# Logistic Regression
lr_model = LogisticRegression()
lr_model.fit(X_train_vec, y_train)
```

```
# Classification report
y_pred = svm_model.predict(X_test_vec)
report = classification_report(y_test, y_pred)print("Classification Report:\n",
report)
```

```
Classification Report:
              precision    recall  f1-score   support

   anger           0.82        0.85        0.83        139
    fear           0.86        0.91        0.88        201
     joy           0.93        0.89        0.91        167
  sadness           0.81        0.74        0.77        122

 accuracy                   0.86                   629
 macro avg           0.85        0.85        0.85        629
weighted avg           0.86        0.86        0.86        629
```

```
# Classification report for Decision Tree
y_pred_dt = dt_model.predict(X_test_vec)
report_dt = classification_report(y_test, y_pred_dt)
print("Decision Tree Classification Report:\n", report_dt)
```

```
Decision Tree Classification Report:
              precision    recall  f1-score   support

   anger         0.80        0.75        0.77        139
    fear         0.74        0.79        0.77        201
     joy         0.86        0.87        0.87        167
  sadness         0.79        0.75        0.77        122

 accuracy                   0.79        629
 macro avg         0.80        0.79        0.79        629
 weighted avg         0.80        0.79        0.79        629
```

```
# Classification report for KNN
y_pred_knn = knn_model.predict(X_test_vec)
report_knn = classification_report(y_test, y_pred_knn)
print("KNN Classification Report:\n", report_knn)
```

```
KNN Classification Report:
              precision    recall  f1-score   support

   anger         0.55        0.74        0.63        139
    fear         0.67        0.69        0.68        201
     joy         0.69        0.57        0.63        167
  sadness         0.56        0.43        0.49        122

 accuracy                   0.62        629
 macro avg         0.62        0.61        0.61        629
 weighted avg         0.63        0.62        0.62        629
```

```
# Classification report for Logistic Regression
y_pred_lr = lr_model.predict(X_test_vec)
report_lr = classification_report(y_test, y_pred_lr)
print("Logistic Regression Classification Report:\n", report_lr)
```

Logistic Regression Classification Report:

	precision	recall	f1-score	support
anger	0.83	0.83	0.83	139
fear	0.81	0.93	0.86	201
joy	0.91	0.89	0.90	167
sadness	0.84	0.66	0.74	122
accuracy			0.84	629
macro avg	0.85	0.83	0.83	629
weighted avg	0.85	0.84	0.84	629

8. Write a program to perform Decision Tree classifier with built-in function on the dataset

```
from google.colab import drive
drive.mount('/content/drive')
```

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

```
data = pd.read_csv('/content/drive/MyDrive/Principles of data science/poo.csv')
data
```

	Unnamed: 0	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2

```
X = data.iloc[:,[1,2]].values
y = data.iloc[:,8].values
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size = 0.25,
random_state= 0)
```

```
#feature scaling
from sklearn.preprocessing import StandardScaler
sc_X = StandardScaler()
X_train = sc_X.fit_transform(X_train)
X_test = sc_X.transform(X_test)
```

```
from sklearn.tree import DecisionTreeClassifier
classifier = DecisionTreeClassifier()
classifier = classifier.fit(X_train,y_train)
```

```
!pip install metrics
```

```
#prediction
y_pred = classifier.predict(X_test)
```

```

#Accuracy
from sklearn import metrics
print('Accuracy Score:', metrics.accuracy_score(y_test,y_pred))

Accuracy score is 0.875

from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)

# Visualizing the Training set results
from matplotlib.colors import ListedColormap

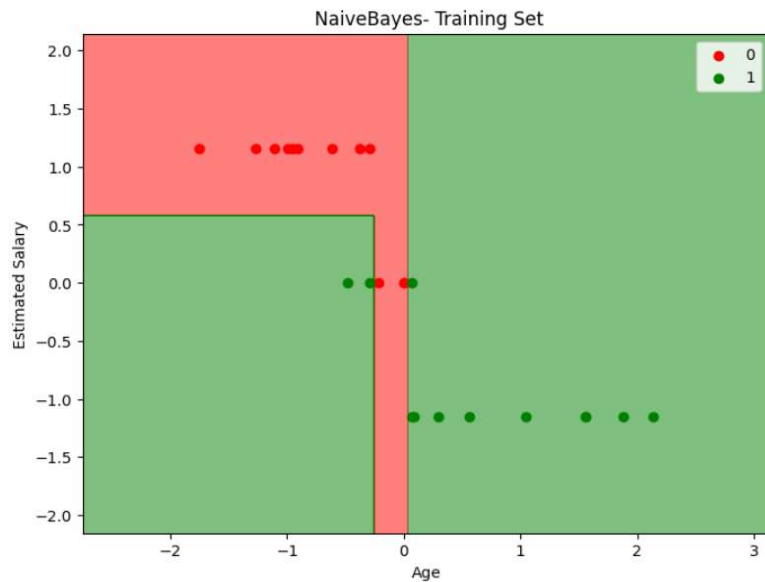
# Create a meshgrid to plot the decision boundary
X1, X2 = np.meshgrid(np.arange(start=X_train[:, 0].min() - 1, stop=X_train[:,
0].max() + 1, step=0.01),
                    np.arange(start=X_train[:, 1].min() - 1, stop=X_train[:, 1].max() + 1,
step=0.01))

# Use the classifier to predict the class labels for each point in the meshgrid
Z = classifier.predict(np.array([X1.ravel(), X2.ravel()]).T)
Z = Z.reshape(X1.shape)

# Create a color map for the plot
cmap = ListedColormap(('red', 'green'))

# Plot the training set data points
plt.figure(figsize=(8, 6))
plt.contourf(X1, X2, Z, alpha=0.5, cmap=cmap)
plt.scatter(X_train[y_train == 0, 0], X_train[y_train == 0, 1], color='red', label='0')
plt.scatter(X_train[y_train == 1, 0], X_train[y_train == 1, 1], color='green', label='1')
plt.title(' NaiveBayes- Training Set')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()

```



!pip install scikit-learn

from sklearn.tree import DecisionTreeClassifier, plot_tree

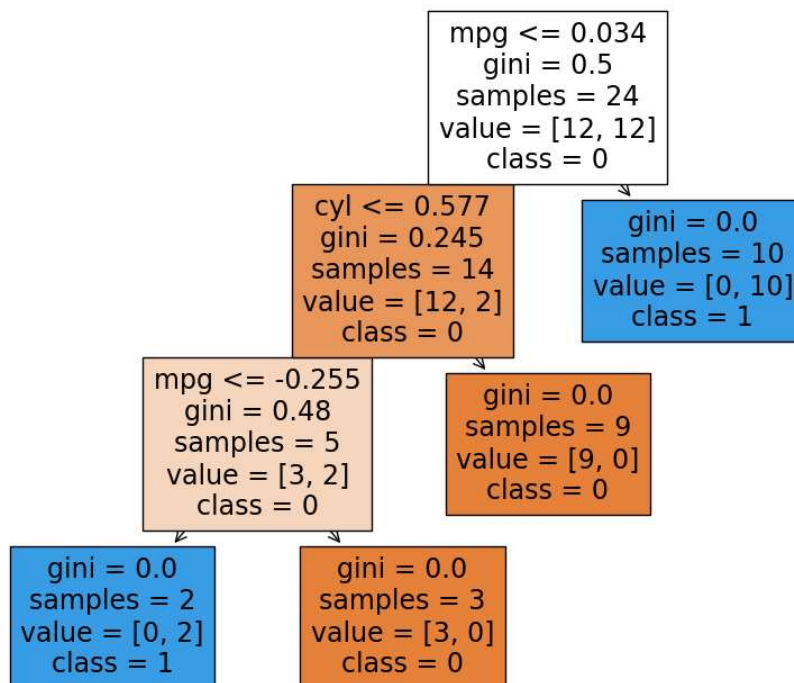
Plotting the Decision Tree graph

plt.figure(figsize=(10, 8))

plot_tree(classifier, feature_names=['mpg', 'cyl'], class_names=['0', '1'],

filled=True)

plt.show()



9. Write a program to perform Density Based Clustering with built-in function on the dataset

```
from numpy import unique
from numpy import where
from matplotlib import pyplot
from sklearn.datasets import make_classification
from sklearn.cluster import DBSCAN
# initialize the data set we'll work with
training_data, _ = make_classification(
    n_samples=1000,
    n_features=2,
    n_informative=2,
    n_redundant=0,
    n_clusters_per_class=1,
    random_state=4
)


# define the model
dbscan_model = DBSCAN(eps=0.25, min_samples=9)

# train the model
dbscan_model.fit(training_data)

# assign each data point to a cluster
dbscan_result = dbscan_model.fit_predict(training_data)

# get all of the unique clusters
dbscan_clusters = unique(dbscan_result)

# plot the DBSCAN clusters
for dbscan_cluster in dbscan_clusters:
    # get data points that fall in this cluster
    index = where(dbscan_result == dbscan_cluster)
    # make the plot
    pyplot.scatter(training_data[index, 0], training_data[index, 1])
    # show the DBSCAN plot
pyplot.show()
```



```
#another method
from sklearn.cluster import DBSCAN
from sklearn.datasets import make_blobs
import matplotlib.pyplot as plt

# Generate sample data with blobs
X, _ = make_blobs(n_samples=500, centers=3, cluster_std=0.6,
random_state=0)

# DBSCAN clustering
db = DBSCAN(eps=0.5, min_samples=5)
y_db = db.fit_predict(X)


# Number of clusters in labels, ignoring noise if present (-1)
n_clusters_ = len(set(y_db)) - (1 if -1 in y_db else 0)

# Plot the clusters
plt.figure(figsize=(8, 6))
plt.scatter(X[y_db == 0][:, 0], X[y_db == 0][:, 1], c='blue', marker='o', label='Cluster 1')
plt.scatter(X[y_db == 1][:, 0], X[y_db == 1][:, 1], c='green', marker='s', label='Cluster 2')
```

```

plt.scatter(X[y_db == 2][:, 0], X[y_db == 2][:, 1], c='red', marker='^', label='Cluster
3')
plt.scatter(X[y_db == -1][:, 0], X[y_db == -1][:, 1], c='gray', marker='x',
label='Noise')
plt.legend(loc='best')
plt.title(f"DBSCAN Clustering (Number of Clusters: {n_clusters_})")
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.show()

```



10. Implement K-Means Clustering with built-in function on the dataset

```
# Importing the libraries
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
# Importing the dataset
dataset = pd.read_csv('content/drive/MyDrive/DATA SCIENCE
LAB/Mall_Customers.csv')
dataset
```

CustomerID Genre Age Annual Income (k\$) Spending Score (1-100)

0 1 Male 19 15 39

198 199 Male 32 137 18


199 200 Male 30 137 83

200 rows × 5 columns

```
import matplotlib.pyplot as plt

# Example data points
spending_scores = dataset.iloc[:,3].values
annual_incomes = dataset.iloc[:, 4].values

# Plotting the scatter plot
plt.scatter(spending_scores,annual_incomes)
plt.title('Spending Score vs. Annual Income')
plt.xlabel('Annual Income')
plt.ylabel('Spending Score')
plt.show()
```



```
X = dataset.iloc[:, [3, 4]].values
```

```
# Fitting K-Means to the dataset
```

```
from sklearn.cluster import KMeans
```

```
kmeans = KMeans(n_clusters = 5, init = 'k-means++', random_state = 42)
```

```
y_kmeans = kmeans.fit_predict(X)
```

```
# Visualising the clusters
```

```
plt.scatter(X[y_kmeans == 0, 0], X[y_kmeans == 0, 1], c = 'red', label = 'Cluster 1')
```

```
plt.scatter(X[y_kmeans == 1, 0], X[y_kmeans == 1, 1], c = 'blue', label = 'Cluster 2')
```

```
plt.scatter(X[y_kmeans == 2, 0], X[y_kmeans == 2, 1], c = 'green', label = 'Cluster 3')
```

```
plt.scatter(X[y_kmeans == 3, 0], X[y_kmeans == 3, 1], c = 'cyan', label = 'Cluster 4')
```

```
plt.scatter(X[y_kmeans == 4, 0], X[y_kmeans == 4, 1], c = 'magenta', label = 'Cluster 5')
```

```
plt.scatter(kmeans.cluster_centers_[0], kmeans.cluster_centers_[1], s = 300, c = 'yellow', label = 'Centroids')
```


```
plt.title('Clusters of customers')
```

```
plt.xlabel('Annual Income (k$)')
```

```
plt.ylabel('Spending Score (1-100)')
```

```
plt.legend()
```

```
plt.show()
```



11. Implement K-Nearest Neighbour classifier with built-in function on the dataset

```
import numpy as nm
import matplotlib.pyplot as mtp
import pandas as pd
```

```
data_set= pd.read_csv('User_Data.csv')
print(data_set)
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
..
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

[400 rows x 5 columns]

```
x= data_set.iloc[:, [2,3]].values
```

```
y= data_set.iloc[:, 4].values
```

```
# Splitting the dataset into training and test set.
```

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test= train_test_split(x, y, test_size= 0.15,
random_state=0)
```

#Feature Scaling

```
from sklearn.preprocessing import StandardScaler
st_x= StandardScaler()
x_train= st_x.fit_transform(x_train)
x_test= st_x.transform(x_test)
```

#Fitting K-NN classifier to the training set

```
from sklearn.neighbors import KNeighborsClassifier
classifier= KNeighborsClassifier(n_neighbors=5, metric='minkowski', p=2 )
classifier.fit(x_train, y_train)
```

#Fitting K-NN classifier to the training set

```
from sklearn.neighbors import KNeighborsClassifier
classifier= KNeighborsClassifier(n_neighbors=5, metric='minkowski', p=2 )
classifier.fit(x_train, y_train)
```

```
▼ KNeighborsClassifier
KNeighborsClassifier()
```

#Predicting the test set result

```
y_pred= classifier.predict(x_test)
print(y_pred)
```

#Creating the Confusion matrix

```
from sklearn.metrics import confusion_matrix
cm= confusion_matrix(y_test, y_pred)
print('Confusion Matrix : \n',cm)
```


Confusion Matrix :

```
[[42  3]
 [ 1 14]]
```

#Visualizing the training set result

from matplotlib.colors import ListedColormap

x_set, y_set = x_train, y_train

x1, x2 = nm.meshgrid(nm.arange(start = x_set[:, 0].min() - 1, stop = x_set[:, 0].max() + 1, step = 0.01),

nm.arange(start = x_set[:, 1].min() - 1, stop = x_set[:, 1].max() + 1, step = 0.01))

mtp.contourf(x1, x2, classifier.predict(nm.array([x1.ravel(),

x2.ravel()]).T).reshape(x1.shape),

alpha = 0.75, cmap = ListedColormap(('red','green')))

mtp.xlim(x1.min(), x1.max())

mtp.ylim(x2.min(), x2.max())

for i, j in enumerate(nm.unique(y_set)):

mtp.scatter(x_set[y_set == j, 0], x_set[y_set == j, 1],

c = ListedColormap(('red', 'green'))(i), label = j)


mtp.title('K-NN Algorithm (Training set)')

mtp.xlabel('Age')

mtp.ylabel('Estimated Salary')

mtp.legend()

mtp.show()



#Visualizing the test set result

from matplotlib.colors import ListedColormap


x_set, y_set = x_test, y_test

x1, x2 = nm.meshgrid(nm.arange(start = x_set[:, 0].min() - 1, stop = x_set[:, 0].max() + 1, step = 0.01),

```

nm.arange(start = x_set[:, 1].min() - 1, stop = x_set[:, 1].max() + 1, step = 0.01))
mtp.contourf(x1, x2, classifier.predict(nm.array([x1.ravel(),
x2.ravel()]).T).reshape(x1.shape),
alpha = 0.75, cmap = ListedColormap(('red','green' )))
mtp.xlim(x1.min(), x1.max())
mtp.ylim(x2.min(), x2.max())
for i, j in enumerate(nm.unique(y_set)):
    mtp.scatter(x_set[y_set == j, 0], x_set[y_set == j, 1],
                c = ListedColormap(('red', 'green'))(i), label = j)
mtp.title('K-NN algorithm(Test set)')
mtp.xlabel('Age')
mtp.ylabel('Estimated Salary')
mtp.legend()
mtp.show()

```



```

from sklearn.metrics import accuracy_score
print ('Accuracy : ', accuracy_score(y_test, y_pred))

```

Accuracy : 0.9333333333333333